

Searching for Reserved Words

An important task of the scanner is to distinguish between user-defined identifiers and reserved words. Since identifiers and reserved words account for a large percentage of the symbols in a typical program, it is important that this task be implemented efficiently. Indeed, an analysis of the collection of correct CPRL test programs shows that almost half the symbols fall into one of these two categories.

The basic idea adopted by many hand-implemented scanners is to accumulate the characters for an identifier or reserved word into a string and then use a look-up mechanism to determine if the string is one of the reserved words. We encapsulate the details of the look-up mechanism in a method with the following signature.

```
public Symbol getIdentifierSymbol(String idString)
```

The method will return either `Symbol.identifier` or one of the reserved word symbols – `Symbol.IntegerRW`, `Symbol.ifRW`, `Symbol.loopRW`, etc.

This handout will explore several algorithms for implementing this method, and it will include benchmark results for the performance of each algorithm. Details of the implementations and benchmark analyses are given in the remainder of this handout, but the results displayed in the following table show that the performance differences can be substantial. Pairwise comparison of the results shows statistically significant differences with the exception of the last two algorithms, the Nested Switch and the HashMap algorithms, which are essentially equivalent.

Search Algorithm	Benchmark Time (in seconds)	Standard Deviation
Sequential 1	21.79	0.9004
Sequential 2	12.36	0.3518
Sequential 3	4.82	0.1271
Binary	5.78	0.1560
By Length	3.87	0.1357
Gperf Hash	3.56	0.1607
By First Character	3.27	0.1224
Nested Switch	2.70	0.1021
HashMap	2.75	0.0994

Benchmarking the Search Algorithms

The general approach to benchmarking was to use each algorithm to perform numerous searches for identifiers and reserved words. The run times for each algorithm were computed simply by reading the system clock before and after performing the searches and then subtracting. While the timings are subject to some noise from other processes running in the background, every effort was made to minimize this effect, and multiple runs of the benchmarks produced comparable times. The results above are obtained by running the timed searches ten times and averaging the results. An initial run of the timed searches was used to “warmup” the JVM, with the results of the initial run discarded.

Since not all reserved words are used with the same frequency, an analysis of the correct CPRL test programs was performed to determine a rough order of magnitude. For example, the reserve word `end` was used a lot more than the reserved word `loop`, and `loop` was used a lot more than the

reserved word `false`. Among predefined type names, `Integer` was by far the most used. Also, recall that some reserved words are not yet used in CPRL but are reserved for possible future use. In addition, user-defined identifiers occurred almost as frequently as all reserved words combined.

Based on this analysis, a file was created to control the number of times each reserved word or identifier would be searched as part of the benchmarking process. Each line of the file contained either a reserved word or an identifier followed by the number of times it was to be searched. Below is an outline of the file.

```
Boolean    1000000
Char       1000000
Integer    4000000
String     500000
and        1000000
array      1000000
begin      4000000
...
type       1000000
var        4000000
when       1000000
while      2000000
write      2000000
writeln    6000000
i          15000000
average    10000000
value1     10000000
x          10000000
thisIsAVeryLongIdentifier 10000000
```

The data in the list was used to initialize a large list of strings, with each identifier or reserved word appearing the specified number of times in the list. After initialization, values in the list were randomly rearranged using the Java method `Collections.shuffle()`, and the list was converted to an array named `testIds` for testing. Timed performance of each search algorithm was implemented as shown below.

```
timer.start();
for (int i = 0; i < testIds.length; ++i)
    search.getIdentifierSymbol(testIds[i]);
timer.stop();
```

Note that this code simply ignores the value returned by method `getIdentifierSymbol()`.

Sequential Search 1

This is the least efficient approach to searching for reserved words and is useful only for comparison purposes. This approach uses an array list of reserved word symbols and a sequential search through the array list. Its only advantage is that it is very easy to implement.

The code to initialize the list is very simple.

```
ArrayList<Symbol> reservedWords = new ArrayList<>();
for (Symbol symbol : Symbol.values())
{
    if (symbol.isReservedWord())
        reservedWords.add(symbol);
}
```

We can visualize the contents of the array list as follows.

```
{
    BooleanRW,
    CharRW,
    IntegerRW,
    StringRW,
    andRW,
    arrayRW,
    beginRW,
    ...
    typeRW,
    varRW,
    whenRW,
    whileRW,
    writeRW,
    writelnRW
}
```

Using this array list, the search method is implemented as shown below.

```
public Symbol getIdentifierSymbol(String idString)
{
    for (int i = 0; i < reservedWords.size(); ++i)
    {
        if (idString.equals(reservedWords.get(i).toString()))
            return reservedWords.get(i);
    }

    return Symbol.identifier;
}
```

Note the call to method `get(i)` for retrieving the i^{th} element in the array list and the call to method `toString()` for retrieving the reserved word spelling.

Sequential Search 2

Similar to the approach above, this approach also uses a sequential search, so we might expect similar performance. But for this approach we make a couple of small changes that greatly improve performance. First, instead of using an array list we use just a plain array. And second, instead of storing simply the symbol and calling `toString()` to get the spelling, we store a pair that has both the spelling (string) and the symbol. We could have used the generic class `Pair` defined in package `javafx.util`, but since JavaFX is no longer shipped with Java and since this is a very simple class, we elected to just implement it directly as a private, static, nested class.

```
private static class StrSymPair
{
    public String rwString;
    public Symbol rwSymbol;

    public StrSymPair(String rwString, Symbol rwSymbol)
    {
        this.rwString = rwString;
        this.rwSymbol = rwSymbol;
    }
}
```

```
}
```

Note that the fields are public to avoid “getX()” method calls, but that is acceptable since class is only used in limited ways by the search methods. In addition, it is declared as private within the classes containing the search algorithm, thereby hiding it from other classes. This class will be also be used by some of the other search methods.

This approach requires only slightly more work initializing the array, and after initialization we can visualize the contents of the array of pairs as follows.

```
{
    ("Boolean", BooleanRW),
    ("Char", CharRW),
    ("Integer", IntegerRW),
    ("String", StringRW),
    ("and", andRW),
    ("array", arrayRW),
    ("begin", beginRW),
    ...
    ("type", typeRW),
    ("var", varRW),
    ("when", whenRW),
    ("while", whileRW),
    ("write", writeRW),
    ("writeln", writelnRW)
}
```

Using this array of (String, Symbol) pairs, the search method is implemented as shown below.

```
public Symbol getIdentifierSymbol(String idString)
{
    for (int i = 0; i < reservedWordPairs.length; ++i)
    {
        if (idString.equals(reservedWordPairs[i].rwString))
            return reservedWordPairs[i].rwSymbol;
    }

    return Symbol.identifier;
}
```

Note that the results shown near the beginning of this handout indicate a run-time performance of roughly half of that required by the **Sequential Search 1** algorithm, a somewhat surprising outcome.

Sequential Search 3

This approach simply uses nested “else if” statements to determine the appropriate identifier symbol.

```
public Symbol getIdentifierSymbol(String idString)
{
    if (idString.equals("Boolean"))
        return Symbol.BooleanRW;
    else if (idString.equals("Char"))
        return Symbol.CharRW;
    else if (idString.equals("Integer"))
        return Symbol.IntegerRW;
```

```

else if (idString.equals("String"))
    return Symbol.StringRW;
else if (idString.equals("and"))
    return Symbol.andRW;
else if (idString.equals("array"))
    return Symbol.arrayRW;
else if (idString.equals("begin"))
    return Symbol.beginRW;
else if (idString.equals("type"))
    ...
    return Symbol.typeRW;
else if (idString.equals("var"))
    return Symbol.varRW;
else if (idString.equals("when"))
    return Symbol.whenRW;
else if (idString.equals("while"))
    return Symbol.whileRW;
else if (idString.equals("write"))
    return Symbol.writeRW;
else if (idString.equals("writeln"))
    return Symbol.writelnRW;
else
    // if you get this far, it must be a plain old identifier
    return Symbol.identifier;
}

```

The performance of this algorithm was surprising. One might expect this algorithm to have roughly the same run-time performance as the **Sequential Search 2** algorithm, but, in fact, it was substantially better.

Binary Search

For this approach we use a binary search of the reserved words with the expectation that it will perform faster. Using big-Oh notation, we know that a sequential search has $O(n)$ performance, while a binary search has $O(\log n)$ performance. It is not difficult to implement a binary search algorithm, but class `Arrays` (in package `java.util`) already does this for us. So first we create an array of strings for the reserved words. We name this array `reservedWordStrs`, and we visualize its contents as follows.

```

{
    "Boolean",
    "Char",
    "Integer",
    "String",
    "and",
    "array",
    "begin",
    ...
    "type",
    "var",
    "when",
    "while",
    "write",

```

```

        "writeln"
    }
}

```

Using `Arrays.binarySearch()` to search this list for string `idString` will return the array index if the string is found; otherwise, it will return a negative value. But when the value is found, we need to convert it to a `Symbol`. There are several ways to do this, but the most efficient way is to have a second (parallel) array of reserved word symbols and then use the index from the string search to retrieve the corresponding symbol. We name this second array `reservedWordSyms`, and we visualize its contents as follows.

```

{
    BooleanRW,
    CharRW,
    IntegerRW,
    StringRW,
    andRW,
    arrayRW,
    beginRW,
    ...
    typeRW,
    varRW,
    whenRW,
    whileRW,
    writeRW,
    writelnRW
}

```

Thus, the search algorithm can be implemented as follows.

```

public Symbol getIdentifierSymbol(String idString)
{
    int index = Arrays.binarySearch(reservedWordStrs, idString);
    return index >= 0 ? reservedWordSyms[index] : Symbol.identifier;
}

```

The benchmark results show that **Binary Search** is more than twice as fast as **Sequential Search 2** but not as fast as **Sequential Search 3**, another somewhat surprising result. Evidently the number of CPRL reserved words is too small to overcome the overhead of implementing a binary search. In other words, we would expect the **Binary Search** to be more efficient if the number of reserved words was much larger. In fact, this was tested and confirmed with Ada reserved words since Ada has approximately twice as many reserved words as CPRL. For Ada reserved words, the performance of the two algorithms was very similar, with **Binary Search** performing slightly better than **Sequential Search 3**. But we can do better.

Search by Length

The general approach here is to use the `StrSymPair` class from **Sequential Search 2**, but, in this case, we put the `(String, Symbol)` pairs in an array of subarrays according to the string length. Then we essentially use the length of the identifier string being searched as a hash to pick out the appropriate subarray, and we then use a sequential search to search the subarray of reserved words having the desired length.

We declare and initialize the array of subarrays as follows.

```

private StrSymPair[][] reservedWordsByLength =

```

```

{
    // first two subarrays (indexes 0 and 1) are empty
    { },
    { },
    {
        // reserved words with two characters
        new StrSymPair("if", Symbol.ifRW),
        new StrSymPair("in", Symbol.inRW),
        new StrSymPair("is", Symbol.isRW),
        new StrSymPair("of", Symbol.ofRW),
        new StrSymPair("or", Symbol.orRW)
    },
    {
        // reserved words with three characters
        new StrSymPair("and", Symbol.andRW),
        new StrSymPair("end", Symbol.endRW),
        new StrSymPair("for", Symbol.forRW),
        new StrSymPair("mod", Symbol.modRW),
        new StrSymPair("not", Symbol.notRW),
        new StrSymPair("var", Symbol.varRW)
    },
    ...

    {
        // reserved words with nine characters
        new StrSymPair("procedure", Symbol.procedureRW),
        new StrSymPair("protected", Symbol.protectedRW)
    }
};

```

Using this array of subarrays, the search algorithm can now be implemented as follows.

```

public Symbol getIdentifierSymbol(String idString)
{
    // quick check based on length
    if (idString.length() > 9)
        return Symbol.identifier;

    // get array of reserved words based on length of idString
    StrSymPair[] reservedWords = reservedWordsByLength[idString.length()];

    // perform a sequential search
    for (int i = 0; i < reservedWords.length; ++i)
    {
        if (idString.equals(reservedWords[i].rwString))
            return reservedWords[i].rwSymbol;
    }

    return Symbol.identifier;
}

```

Benchmark results indicate that this search algorithm performs faster than any of the sequential searches or **Binary Search**. It is possible that we could tweak out a few extra milliseconds by reordering the subarrays according to expected frequency of use (e.g., moving `Symbol.loopRW` to the

beginning of the subarray for reserved words of length four and moving `Symbol.trueRW` to the end). Not shown are the results of using a binary search within the appropriate subarray, but the subarrays are small enough that using a binary search resulted in poorer performance. Besides, we can do better.

Search Using Gperf Hash

The need to perform a fast search for a list of words, reserved or otherwise, has been around since the early days of computing, so it should come as no surprise that there is a software utility that can provide assistance. Gperf is a perfect hash function generator for a list of strings. You give gperf a list of strings, and it produces a hash function and hash table in the form of C/C++ code for looking up an arbitrary string to see if it is one of those in the list. The hash function is perfect meaning that the hash table has no collisions and the lookup needs only one single string comparison. Gperf uses extensive analysis of the list of strings (plus a little magic and trial and error) to generate the C/C++ code. Gperf is used for keyword recognition in several production and research compilers including GNU C and GNU C++.

Gperf is available primarily on Linux, but there are versions available for Windows or through Windows Subsystem for Linux (WSL). Gperf on Ubuntu (running on WSL) was used to generate the C code for CPRL reserved words, and the C code was translated to Java for comparison purposes. The output of gperf is a little cryptic. It has three major components.

First, we have `hashCode`, an array that maps characters to integer values. Here is the declaration for this array.

```
private int[] hashCode = new int[256];
```

Most of the positions in the array have value 70, but characters that appeared frequently in CPRL reserved words had different values. Here is an outline of the code used to initialize the `hashCode`.

```
for (int i = 0; i < hashCode.length; ++i)
    hashCode[i] = 70;

hashCode[(int)'B'] = 40;
...
hashCode[(int)'e'] = 5;
hashCode[(int)'f'] = 0;
...
hashCode[(int)'i'] = 5;
hashCode[(int)'l'] = 5;
hashCode[(int)'m'] = 25;
hashCode[(int)'n'] = 5;
hashCode[(int)'o'] = 0;
hashCode[(int)'p'] = 5;
...
hashCode[(int)'y'] = 0;
```

The idea is that the letter 'B' is associated with the integer 40 using this array, and the letter 'e' is associated with the integer 5. The selection of integer values is certainly not obvious at first glance. (I told you it was magic.)

Second, we associate a hash value (an integer) with a candidate string by adding the string's length, the hash code for the character at position zero, the hash code for the character at position one, and the hash code at position three. Assuming that `str` is the name of a candidate string, we compute a hash value for `str` as follows.


```
hash(str) = str.length + hashCode[str.charAt(0)] + hashCode[str.charAt(1)]
           + hashCode[str.charAt(3)]
```

(More magic!) If the string is short, we omit terms for those character positions. Let's look at a few examples.

The hash value for reserved word "Boolean" is given by $7 + 40 + 0 + 5 = 52$.

The hash value for reserved word "loop" is given by $4 + 5 + 0 + 5 = 14$.

The hash value for reserved word "of" is given by $2 + 0 + 0 = 2$.

The hash value for user defined identifier "john" is given by $4 + 70 + 0 + 5 = 79$.

(Note that the letter 'j' has a default hashCode value of 70.)

Third, we create an array of symbols that are indexed by hash values for our reserved words. Since the hash function isn't minimal, there will be entries in the array that do not correspond to a reserved word. For those entries we use the value `Symbol.unknown`. Our array contains 70 symbols and is defined as follows.

```
private Symbol[] symbolList =
{
    Symbol.unknown,      Symbol.unknown,      Symbol.ofRW,
    Symbol.forRW,        Symbol.unknown,      Symbol.unknown,
    Symbol.unknown,      Symbol.ifRW,         Symbol.notRW,
    Symbol.typeRW,       Symbol.unknown,      Symbol.StringRW,
    Symbol.inRW,         Symbol.endRW,        Symbol.loopRW,
    Symbol.beginRW,      Symbol.publicRW,     Symbol.orRW,
    Symbol.andRW,        Symbol.elseRW,       Symbol.elseifRW,
    Symbol.unknown,      Symbol.isRW,         Symbol.varRW,
    Symbol.trueRW,       Symbol.writeRW,      Symbol.returnRW,
    Symbol.writelnRW,    Symbol.modRW,        Symbol.protectedRW,
    Symbol.falseRW,      Symbol.unknown,      Symbol.declareRW,
    Symbol.functionRW,   Symbol.exitRW,       Symbol.unknown,
    Symbol.unknown,      Symbol.privateRW,    Symbol.unknown,
    Symbol.readRW,       Symbol.arrayRW,      Symbol.readlnRW,
    Symbol.programRW,    Symbol.unknown,      Symbol.unknown,
    Symbol.constRW,      Symbol.unknown,      Symbol.IntegerRW,
    Symbol.unknown,      Symbol.unknown,      Symbol.classRW,
    Symbol.unknown,      Symbol.BooleanRW,    Symbol.unknown,
    Symbol.procedureRW,  Symbol.unknown,      Symbol.unknown,
    Symbol.unknown,      Symbol.unknown,      Symbol.thenRW,
    Symbol.unknown,      Symbol.unknown,      Symbol.unknown,
    Symbol.unknown,      Symbol.whenRW,       Symbol.whileRW,
    Symbol.unknown,      Symbol.unknown,      Symbol.unknown,
    Symbol.CharRW
};
```

(Yet more magic!) Note that, as expected, `Symbol.BooleanRW` is at index 52, `Symbol.loopRW` is at index 14, and `Symbol.ofRW` is at index 2.

With these three components in place, our search algorithm is implemented as follows.

```
public Symbol getIdentifierSymbol(String idString)
{
    if (idString.length() >= 2 && idString.length() <= 9)
    {
```

```

        int key = hash(idString);

        if (key < symbolList.length)
        {
            Symbol s = symbolList[key];

            if (idString.equals(s.toString()))
                return s;
        }

        return Symbol.identifier;
    }
}

```

Benchmark results indicate that this search algorithm is better than the **By Length** search shown above and slightly slower than the **Search by First Character** shown below.

Search by First Character

This approach is somewhat similar to that of **Search by Length**, but, rather than organizing the array of subarrays according to string length, we organize according to the first character in the string. Then we use the first character of the identifier string being searched as a hash to pick out the appropriate subarray. The primary advantage here is that the subarrays are much shorter (average 2.1 items per nonempty subarray versus 5.1 for **Search by Length**), so sequentially searching a subarray takes less time. The primary disadvantage is that it is a little more complicated to initialize the arrays.

We declare and initialize the array of subarrays as follows.

```

private StrSymPair[][] words =
{
    // first 65 subarrays (indexes 0-64) are empty
    { }, { }, { }, { }, { }, { }, { }, { }, { }, { }, { },
    ...
    { }, { }, { }, { }, { }, { }, { }, { }, { }, { }, { },
    { }, { }, { }, { }, { },
    { },    // 'A'
    { new StrSymPair("Boolean",    Symbol.BooleanRW) },
    { new StrSymPair("Char",        Symbol.CharRW) },
    { },    // 'D'
    ...
    { },    // 'H'
    { new StrSymPair("Integer",     Symbol.IntegerRW) },
    { },    // 'J'
    ...
    { },    // 'R'
    { new StrSymPair("String",      Symbol.StringRW) },
    { },    // 'T'
    ...
    { },    // ``
    { new StrSymPair("and",          Symbol.andRW),
      new StrSymPair("array",       Symbol.arrayRW) },
    { new StrSymPair("begin",       Symbol.beginRW) },
    { new StrSymPair("class",       Symbol.classRW),

```

```

        new StrSymbPair("const",      Symbol.constRW) },
    { new StrSymbPair("declare",      Symbol.declareRW) },
    { new StrSymbPair("else",         Symbol.elseRW),
      new StrSymbPair("elsif",       Symbol.elsifRW),
      new StrSymbPair("end",          Symbol.endRW),
      new StrSymbPair("exit",         Symbol.exitRW) },
    ...
    { },    // 's'
    { new StrSymbPair("then",          Symbol.thenRW),
      new StrSymbPair("true",          Symbol.trueRW),
      new StrSymbPair("type",          Symbol.typeRW) },
    { },    // 'u'
    { new StrSymbPair("var",            Symbol.varRW) },
    { new StrSymbPair("when",           Symbol.whenRW),
      new StrSymbPair("while",         Symbol.whileRW),
      new StrSymbPair("write",         Symbol.writeRW),
      new StrSymbPair("writeln",       Symbol.writelnRW) },
    { },    // 'x'
    { },    // 'y'
    { }     // 'z'
};

```

Using this array of subarrays, the search algorithm can now be implemented as follows.

```

public Symbol getIdentifierSymbol(String idString)
{
    // get array of reserved word pairs based on first char of idString
    StrSymbPair[] reservedWordPairs = words[(int) idString.charAt(0)];

    // perform a sequential search
    for (StrSymbPair rwPair : reservedWordPairs)
    {
        if (idString.equals(rwPair.rwString))
            return rwPair.rwSymbol;
    }

    return Symbol.identifier;
}

```

Benchmark results show that this algorithm is, indeed, faster than searching based on the length of the string, and it is slightly faster than the **Gperf Hash** search.

Nested Switch

This algorithm, while very efficient, is a little more complicated to implement. We start with a quick check to eliminate strings of length 1 followed by a switch based on the first character (the character at position 0). Here is an outline that shows these first steps with the “outer” switch. Note that we need to include cases only for those letters that can start a reserved word. For example, if parameter `idString` starts with the character 'A' or the character 'g', then `idString` cannot be a reserved word, and the return value from the switch is handled by the default case.

```

public Symbol getIdentifierSymbol(String idString)
{
    // quick check based on length

```

```

    if (idString.length() < 2)
        return Symbol.identifier;

    switch(idString.charAt(0))
    {
        case 'B':    ...
        case 'C':    ...
        case 'I':    ...
        case 'S':    ...
        case 'a':    ...
        case 'b':    ...
        case 'c':    ...
        case 'd':    ...
        case 'e':    ...
        case 'f':    ...
        case 'i':    ...
        case 'l':    ...
        case 'm':    ...
        case 'n':    ...
        case 'o':    ...
        case 'p':    ...
        case 'r':    ...
        case 'v':    ...
        case 't':    ...
        case 'w':    ...
        default:
            return Symbol.identifier;
    }
}

```

We handle each case separately depending on the reserved words that can start with that character. In some cases, there is only one possible reserved word that starts with that character, so the code for that case is simple. For example, the only reserved word that starts with the letter 'B' is Boolean, so that case is implemented as follows.

```

    return idString.equals("Boolean") ? Symbol.BooleanRW : Symbol.identifier;

```

Letters 'C' (Char), 'I' (Integer), 'S' (String), 'b' (begin), 'd' (declare), 'l' (loop), 'm' (mod), 'n' (not), and 'v' (var) are handled similarly.

Most other cases can be handled by a nested switch statement. For example, two reserved words start with the letter 'a' (and, array), and, since these two reserved words differ in the second character (the character at position 1), they are handled as shown.

```

    switch(idString.charAt(1))
    {
        case 'n':
            return idString.equals("and") ? Symbol.andRW : Symbol.identifier;
        case 'r':
            return idString.equals("array") ? Symbol.arrayRW : Symbol.identifier;
        default:
            return Symbol.identifier;
    }
}

```

Similarly, three reserved words start with the letter 'f'.

```

switch(idString.charAt(1))
{
    case 'u':
        return idString.equals("function") ?
            Symbol.functionRW : Symbol.identifier;
    case 'a':
        return idString.equals("false") ? Symbol.falseRW : Symbol.identifier;
    case 'o':
        return idString.equals("for") ? Symbol.forRW : Symbol.identifier;
    default:
        return Symbol.identifier;
}

```

The case where the first letter is 'e' (else, elsif, end, exit) is handled similarly but with a slight twist since both else and elsif have the same character at position 1.

```

switch(idString.charAt(1))
{
    case 'l':
        {
            if (idString.equals("else"))
                return Symbol.elseRW;
            else if (idString.equals("elsif"))
                return Symbol.elsifRW;
            else
                return Symbol.identifier;
        }
    case 'n':
        return idString.equals("end") ? Symbol.endRW : Symbol.identifier;
    case 'x':
        return idString.equals("exit") ? Symbol.exitRW : Symbol.identifier;
    default:
        return Symbol.identifier;
}

```

We show one other interesting case. There are five reserved words that begin with the letter 'p' (private, procedure, program, protected, public), but using a nested switch based on the character at position 1 is not much help since 4 of the 5 reserved words have the same character at that position. Instead, we use a nested switch based on the length of the string.

```

switch(idString.length())
{
    case 6:
        return idString.equals("public") ? Symbol.publicRW : Symbol.identifier;
    case 7:
        {
            if (idString.equals("program"))
                return Symbol.programRW;
            else if (idString.equals("private"))
                return Symbol.privateRW;
            else
                return Symbol.identifier;
        }
    case 9:
        {

```

```

        if (idString.equals("procedure"))
            return Symbol.procedureRW;
        else if (idString.equals("protected"))
            return Symbol.protectedRW;
        else
            return Symbol.identifier;
    }
    default:
        return Symbol.identifier;
}

```

While this algorithm is essentially tied for the fastest, it is also the longest by far. The number of lines of code used to implement this algorithm is about twice as long as any of the other algorithms, and it is about three times as long as the average of the other algorithms. The next algorithm is just as fast, and its implementation requires only a fraction of the number of lines of code.

Search Using HashMap

Our last search uses class `HashMap` from package `java.util`. `HashMap` maps keys to values, and, if the key type has an efficient `hashCode()` method, `HashMap` can be very fast. In our case we want to map objects of type `String` to object of type `Symbol`, so we declare our map as follows.

```
private HashMap<String, Symbol> rwMap;
```

Fortunately for us, `String` is predefined in Java, and it has a very efficient `hashCode()` method.

The remaining details are straightforward and are left as an exercise, but initialization of the `HashMap` is just as simple and straightforward as it was to initialize the `ArrayList` in **Sequential Search 1**, and, based on the benchmark analysis, using **HashMap** is essentially tied with **Nested Switch** for the fastest search algorithm.

Easy to implement and fastest performance, it is the best of both worlds.